

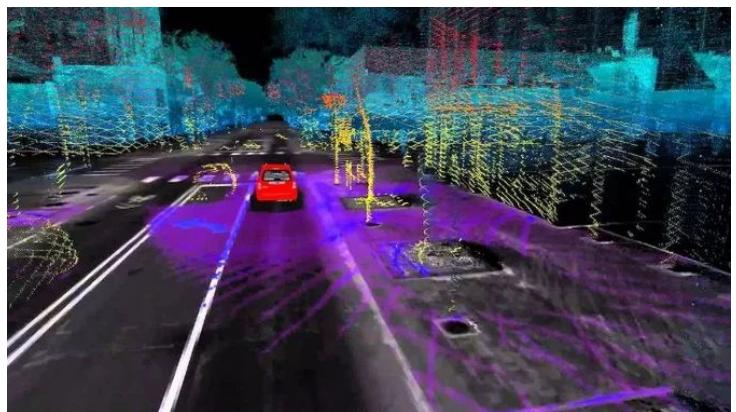
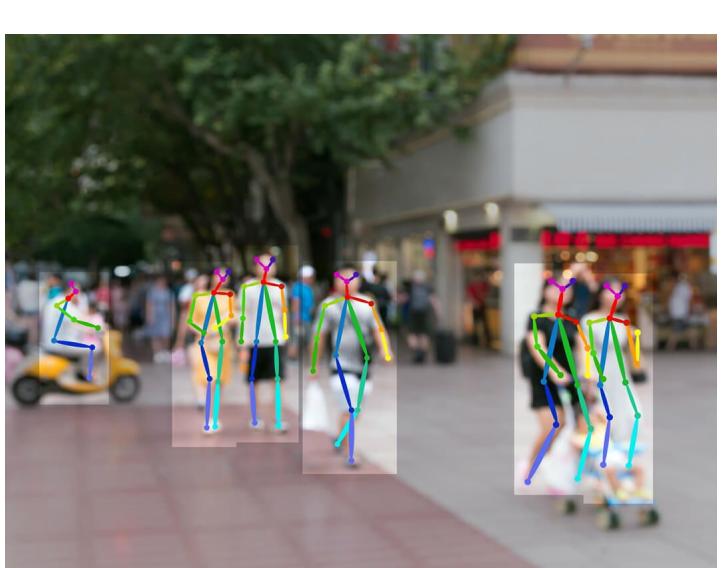
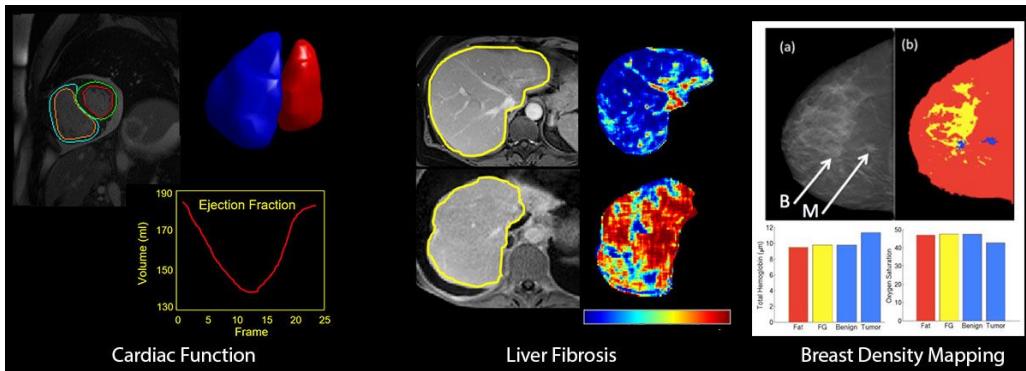
# Image processing

Rotem Shalev-Arkushin  
rotems7@mail.tau.ac.il

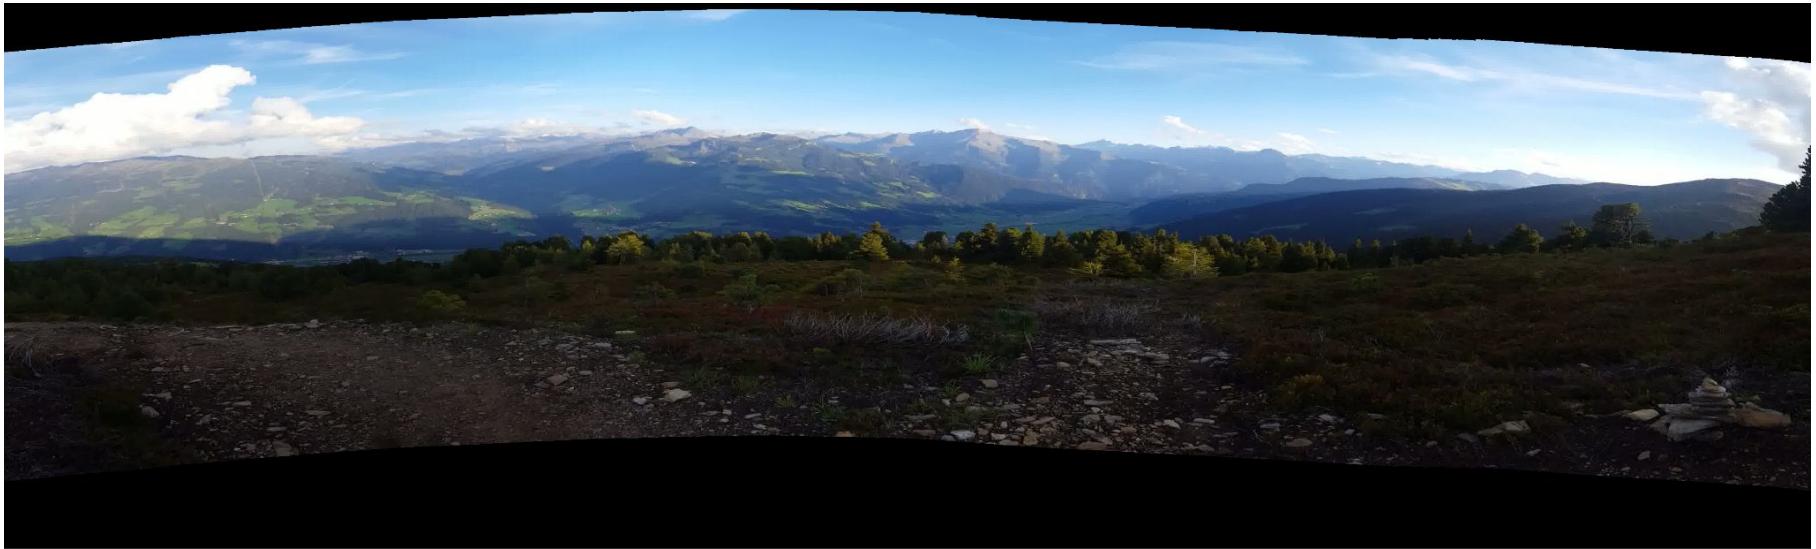
December 17 2025

# Applications

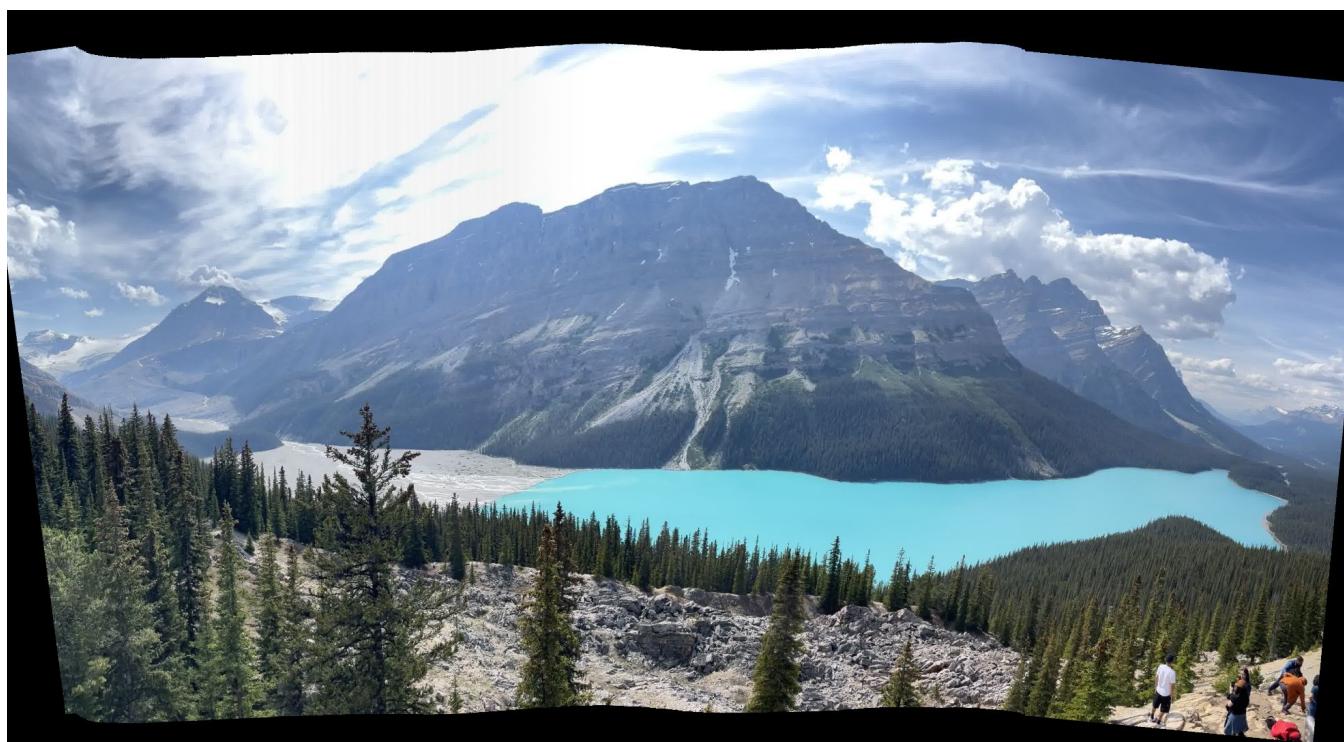
- **Image Enhancement** – deblurring, denoising, sharpening, etc.
- **Registration & Alignment** – panorama stitching, satellite image alignment, AR overlay.
- **Segmentation** – used for object detection and recognition, e.g. for medical imaging.
- **Measurement & Inspection** – detecting defects in manufacturing, counting cells, measuring distances or motion.
- **Tracking & Motion Analysis** – optical flow, gesture tracking, traffic monitoring, video stabilization.
- **3D Reconstruction** – structure-from-motion, depth estimation, multi-view geometry, robotics navigation.
- **Compression** – pyramids, transforms, and feature extraction used in JPEG, wavelets, video codecs.
- **Augmented Reality & Robotics** – feature matching, localization, SLAM, obstacle detection.



# Panorama



# Panorama exercise



# Panorama exercise



# Panorama exercise

What do we need to know to construct a panorama from video frames?



# Panorama exercise

What do we need to know to construct a panorama?

- Image registration and alignment

- Find unique features in each image (use feature detectors and descriptors)
- Find matching features between images
- Align images

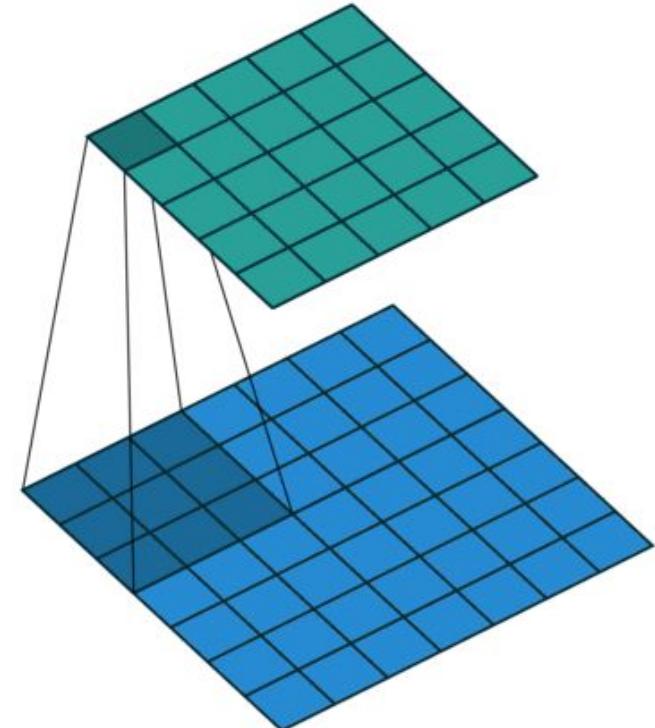
- Stitch images together

- Place aligned images on a shared canvas
- Blend images together seamlessly

## Reminder - 2D convolution

$$g(x, y) = \omega * f(x, y) =$$

$$\sum_{i=-a}^a \sum_{j=-b}^b \omega(i, j)f(x - i, y - j)$$



\* Illustration from Aaditya Prakash's [blog](#)

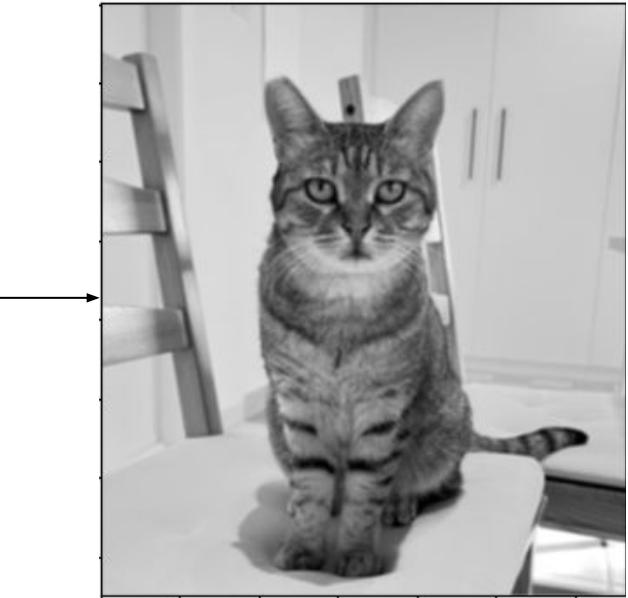
# Image Blurring/smoothing

Convolve with a blurring kernel.

In Python: `convolve2d(koshka_im_gray, kernel, mode='same', boundary='symm')`

- Box blurring

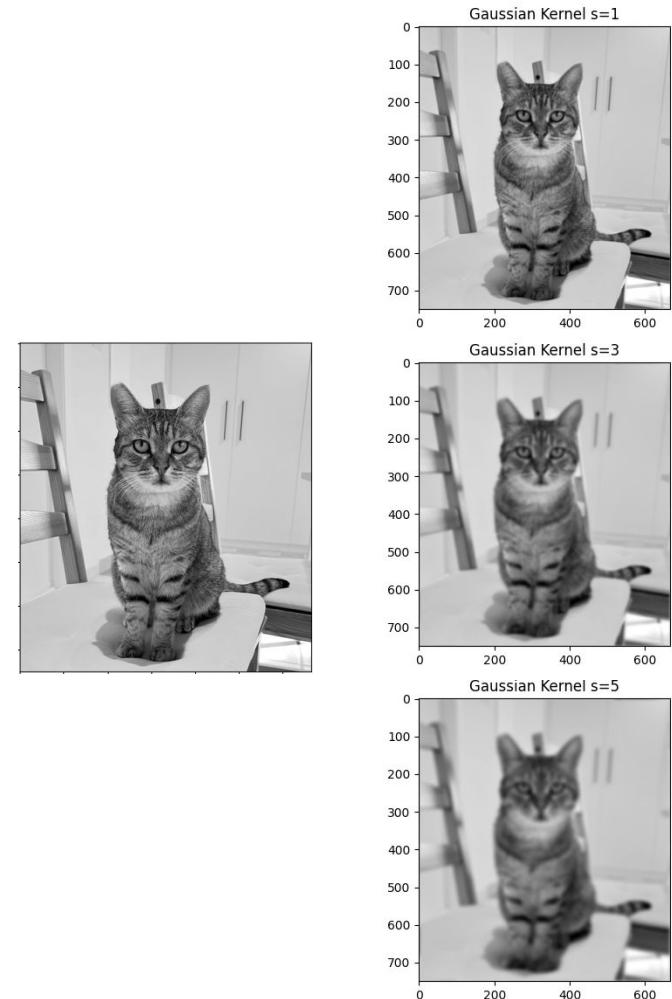
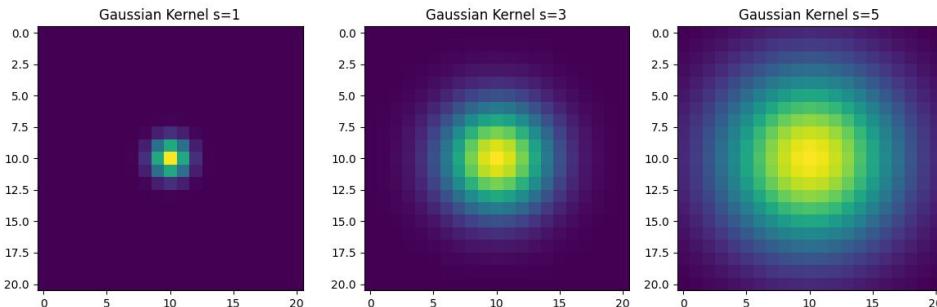
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



# Image Blurring/smoothing

- Convolve with a Gaussian kernel

$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



# Pascal's Triangle

- An infinite triangular array of the binomial coefficients:

$$\binom{n}{k} = \frac{n!}{k!(n-k)!} = \binom{n-1}{k-1} + \binom{n-1}{k}$$

- De Moivre–Laplace theorem:

$$\frac{1}{2^n} \binom{n}{k} \lim_{n \rightarrow \infty} \frac{1}{2\pi\sigma^2} e^{-\frac{((k-n)/2)^2}{2\sigma^2}}$$

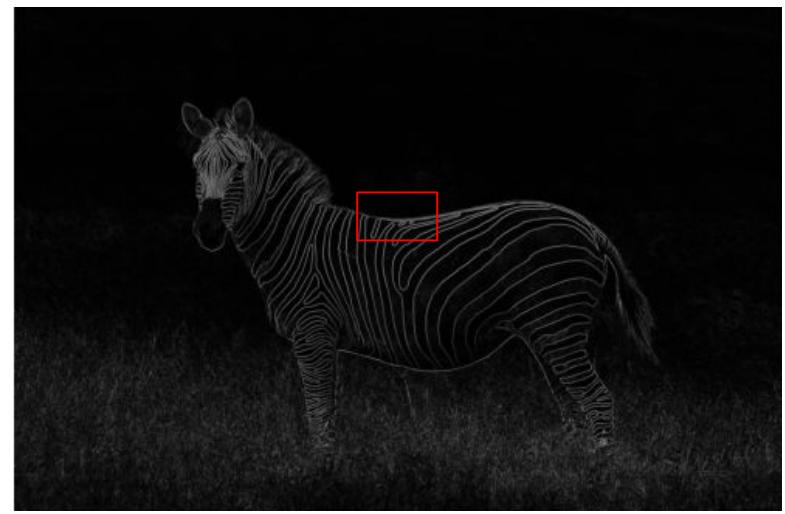
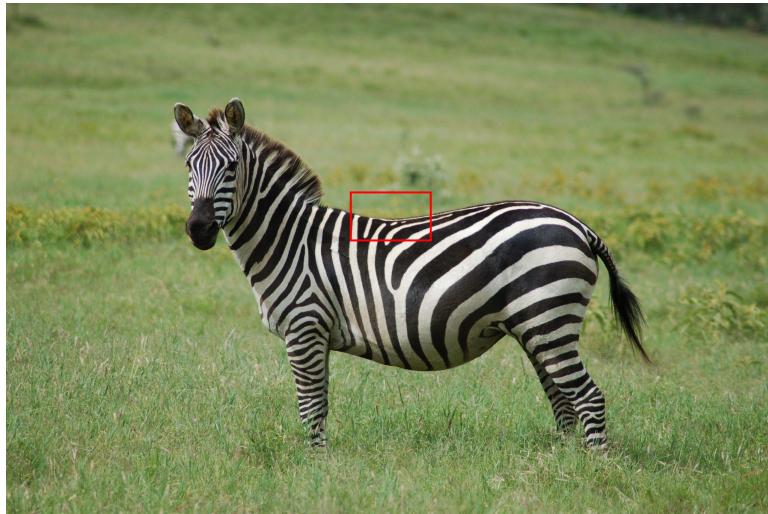
- Repeated convolution with [1 1] generates a sequence of kernels whose shape approaches a Gaussian => binomial filters act as discrete Gaussian approximations.

$$\begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

# Edge detection

- Edge in an image = a line of high intensity change



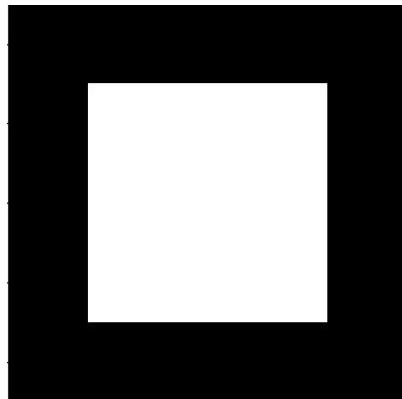
# Edge detection

- Edge in an image = a line of high intensity change
- Why is it interesting?  
segmentation, feature detection, image alignment, recognizing shapes or structures in scenes, controlled generation, and more.
- Gradient (vector of the derivatives) = direction of most rapid intensity change.

$$\frac{d}{dx} f(i, j) = \lim_{\epsilon \rightarrow 0} \frac{f(i, j) - f(i - \epsilon, j)}{\epsilon} \simeq f(i, j) - f(i - 1, j)$$

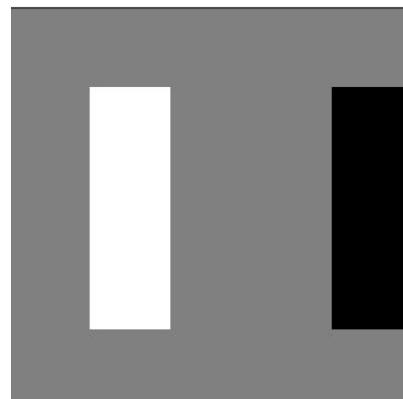
# Image Gradient

0	0	0	0	0
0	255	255	255	0
0	255	255	255	0
0	255	255	255	0
0	0	0	0	0

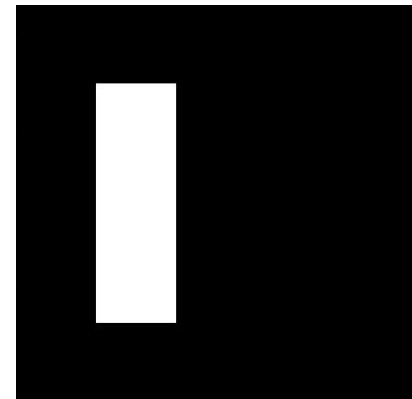


\* [1, -1] =

0	0	0	0	0
0	255	0	0	-255
0	255	0	0	-255
0	255	0	0	-255
0	0	0	0	0



→ clip(0,255)

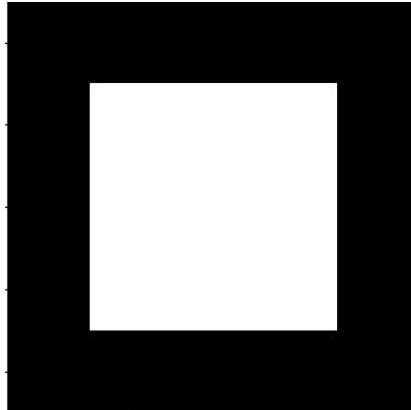


# Image Gradient

0	0	0	0	0
0	255	255	255	0
0	255	255	255	0
0	255	255	255	0
0	0	0	0	0

$$* [1, -1]^T =$$

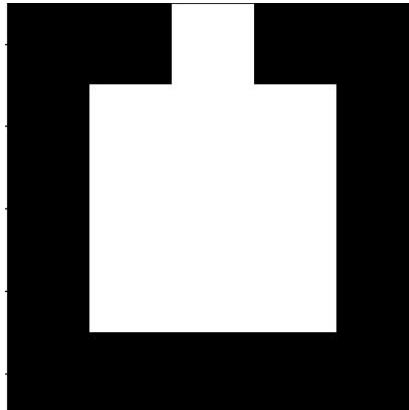
0	0	0	0	0
0	255	255	255	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0



# Image Gradient

0	0	<b>255</b>	0	0
0	255	255	255	0
0	255	255	255	0
0	255	255	255	0
0	0	0	0	0

\*  $[1, -1]^T =$

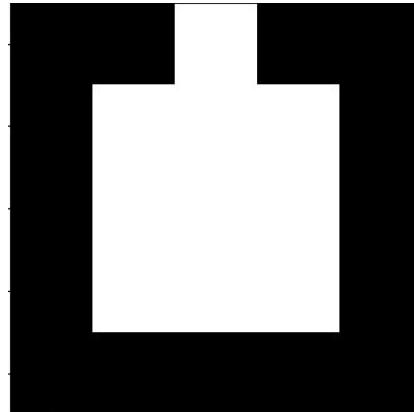


Boundaries?

- Zero-padding
- Mirror last
- Cyclic? **No!**

# Image Gradient

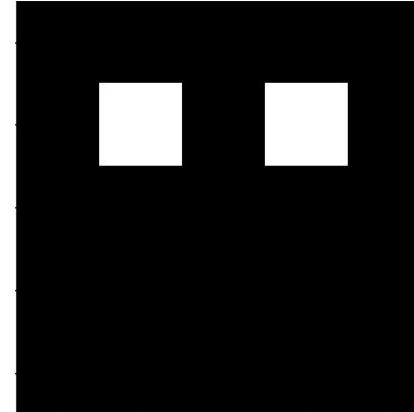
0	0	255	0	0
0	255	255	255	0
0	255	255	255	0
0	255	255	255	0
0	0	0	0	0



Symmetric / mirror

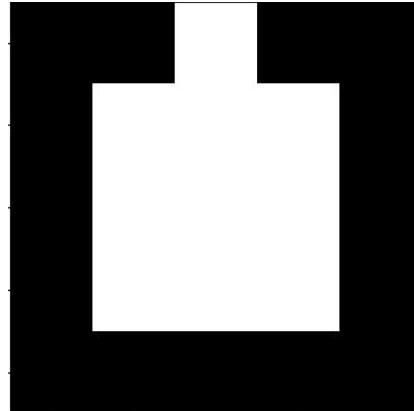
0	0	0	0	0
0	255	0	255	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$$* [1, -1]^T =$$



# Image Gradient

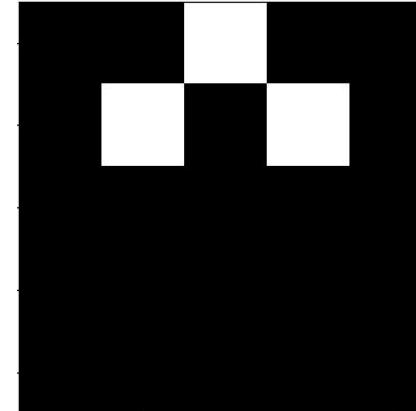
0	0	255	0	0
0	255	255	255	0
0	255	255	255	0
0	255	255	255	0
0	0	0	0	0



zero-pad / cyclic (wrap)

0	0	255	0	0
0	255	0	255	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$$* [1, -1]^T =$$

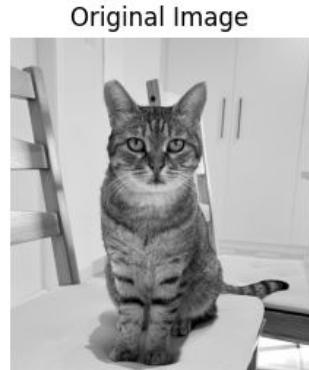


# Derivative Approximation

$$f'(x) \approx \frac{f(x+1) - f(x-1)}{2}$$

With Convolution:

- use kernel of [1 0 -1] and detect vertical edges in an image.
- use kernel of  $[1 \ 0 \ -1]^T$  and detect horizontal edges in an image.



# Sobel

- To avoid noise – smooth the image and then apply derivative.

$$(Im * \frac{1}{4} [1 \ 2 \ 1]) * [-1 \ 0 \ 1]$$

- Sobel – apply them both together with one convolution!

$$S_x = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} [-1 \ 0 \ 1] = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$S_y = \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} [1 \ 2 \ 1] = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

# Sobel

Gradient magnitude and direction:

$$|\nabla f| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2} \quad \alpha = \tan^{-1}\left(\left(\frac{\partial f}{\partial y}\right) / \left(\frac{\partial f}{\partial x}\right)\right)$$

Original Image



Vertical Edges (Sobel)



Horizontal Edges (Sobel)



Gradient Magnitude



# Canny Edge Detector

## Algorithm:

1. Smooth the image with a Gaussian filter to remove noise.
2. Find the intensity gradients of the image with a simple kernel (e.g.  $\frac{1}{2} [1, 0, -1]$ ):  $G_x$ ,  $G_y$ .
3. Calculate the gradient direction:  $\tan(\alpha) = \frac{G_x}{G_y}$  and magnitude:  $\sqrt{G_x^2 + G_y^2}$
4. **Non-maximum suppression**: only keep local magnitude maxima in each edge direction.

# Non-maximum suppression

Before NMS



After NMS



# Non-maximum suppression

**Angles**

18	96	103	100	0.08
15	95	100	97	12
0.72	102	125	124	26
0.4	102	125	125	70
18	96	103	101	0.08

**Magnitudes**

26	45	78	18	18
29	43	80	42	13
30	36	78	47	8
32	30	76	51	4
35	25	75	54	2

# Non-maximum suppression

**Angles**

18	96	103	100	0.08
15	95	100	97	12
0.72	102	125	124	26
0.4	102	125	125	70
18	96	103	101	0.08

**Magnitudes**

0	45	78	18	18
0	43	80	42	0
0	36	78	47	8
0	30	76	51	4
35	25	75	54	2

$0 \leq a \leq 22$

# Non-maximum suppression

**Angles**

18	96	103	100	0.08
15	95	100	97	12
0.72	102	125	124	26
0.4	102	125	125	70
18	96	103	101	0.08

**Magnitudes**

0	0	78	0	18
0	0	80	0	0
0	36	78	47	8
0	30	76	51	4
35	0	75	0	2

$67.5 \leq a \leq 112.5$

# Non-maximum suppression

**Angles**

18	96	103	100	0.08
15	95	100	97	12
0.72	102	125	124	26
0.4	102	125	125	70
18	96	103	101	0.08

**Magnitudes**

0	0	78	0	18
0	0	80	0	0
0	36	78	47	8
0	30	76	0	4
35	0	75	0	2

$112.5 \leq a \leq 157.5$

# Non-maximum suppression

Before NMS

26	45	78	18	18
29	43	80	42	13
30	36	78	47	8
32	30	76	51	4
35	25	75	54	2

After NMS

0	0	<b>78</b>	0	<b>18</b>
0	0	<b>80</b>	0	0
0	<b>36</b>	<b>78</b>	<b>47</b>	<b>8</b>
0	<b>30</b>	<b>76</b>	0	<b>4</b>
<b>35</b>	0	<b>75</b>	0	2

# Canny Edge Detector

## Algorithm:

1. Smooth the image with a Gaussian filter to remove noise.
2. Find the intensity gradients of the image with a simple kernel (e.g.  $\frac{1}{2} [1, 0, -1]$ ):  $G_x$ ,  $G_y$ .
3. Calculate the gradient direction:  $\tan(\alpha) = \frac{G_x}{G_y}$  and magnitude:  $\sqrt{G_x^2 + G_y^2}$
4. **Non-maximum suppression**: only keep local magnitude maxima in each edge direction.
5. **Hysteresis with double threshold ( $t_1 > t_2$ )**: keep edges  $> t_1$ .  
Link  $t_1 >$  edges  $> t_2$  – keep only edges neighboring an edge  $> t_1$ .

# Hysteresis

**Before Hysteresis**

0	0	78	0	18
0	0	80	0	0
0	36	78	47	8
0	30	76	0	4
<b>35</b>	0	75	0	2

**After Hysteresis**

0	0	78	0	0
0	0	80	0	0
0	0	78	47	0
0	0	76	0	0
0	0	75	0	0

$$t1 = 77, t2 = 40$$

# Canny Edge Detector

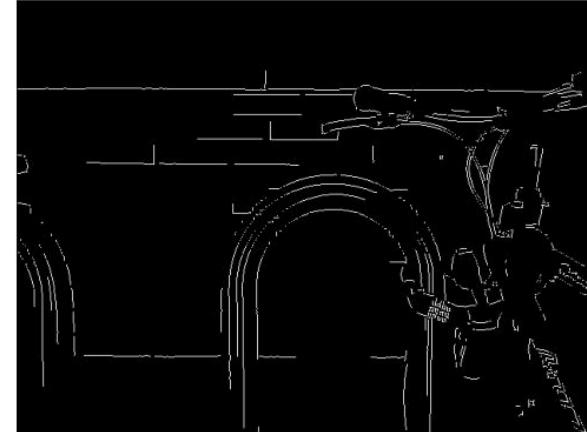
Initial magnitude

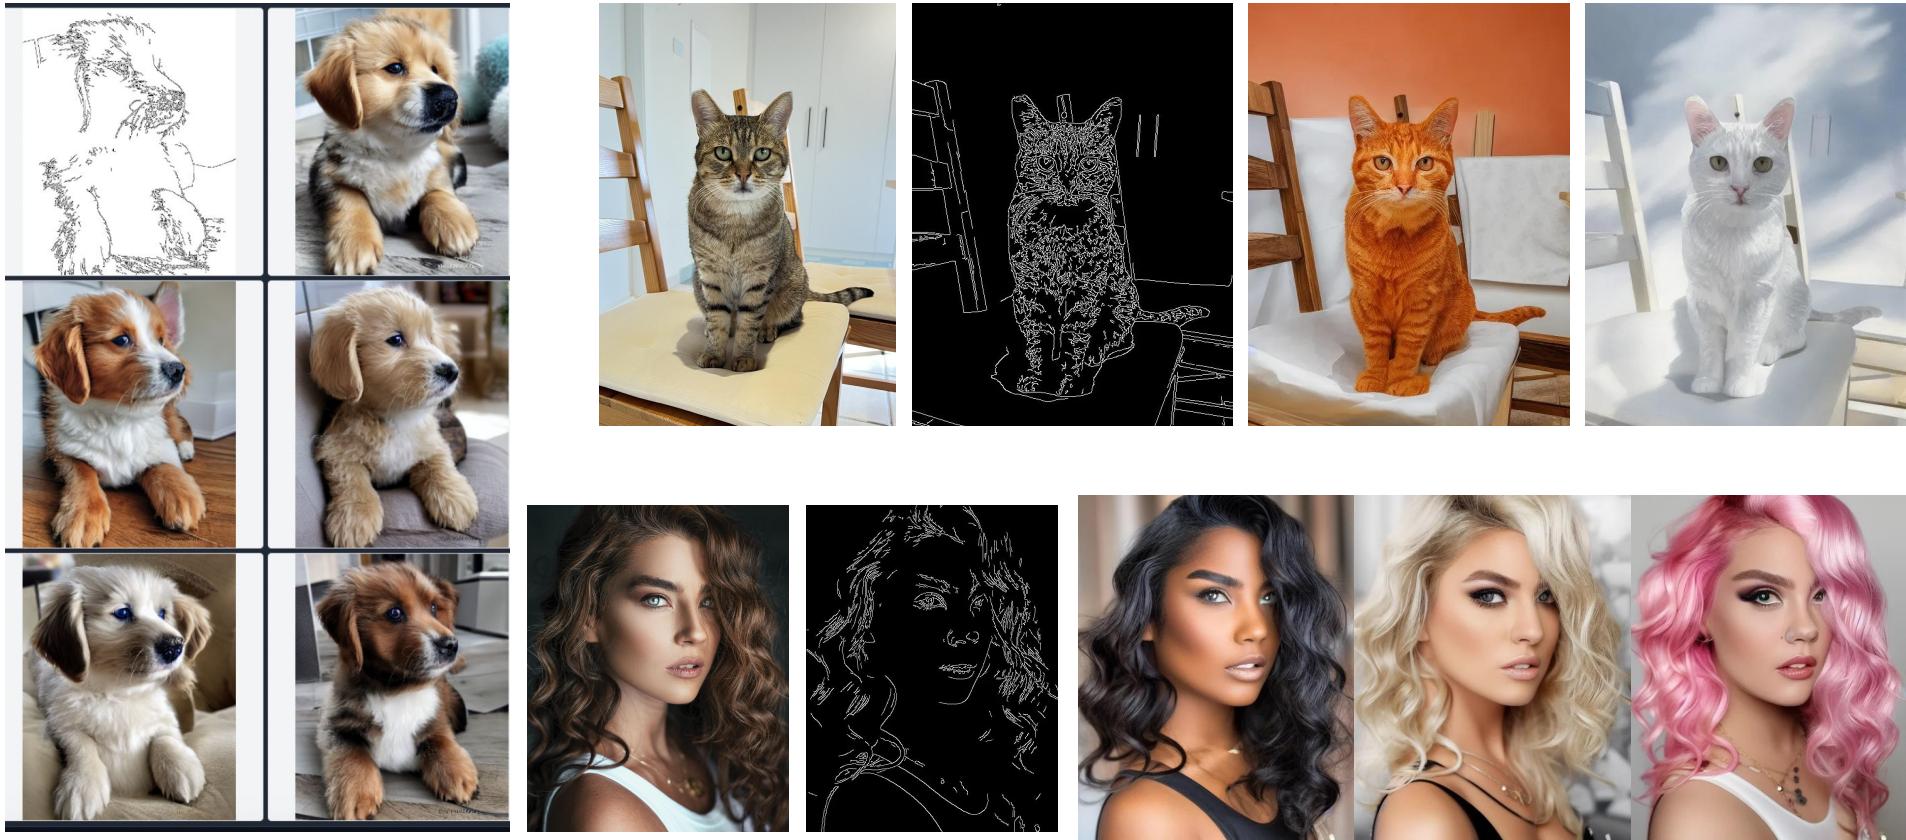


$> t_1$



After hysteresis





ControlNet generations with Canny edges

# Binary images

- Images that are only black and white (pixel values are either 0 or 255).
- Used for segmentation, masking for inpainting, controlled generation, etc.



# Pixel Neighborhoods and Connectivity

- **Connected Component:** A set of pixels is a Connected Component if for every two pixels in the set, there is a path between them such that every two successive pixels in the path are in the set, and they are neighbors (could be with 4 or 8 connectivity).

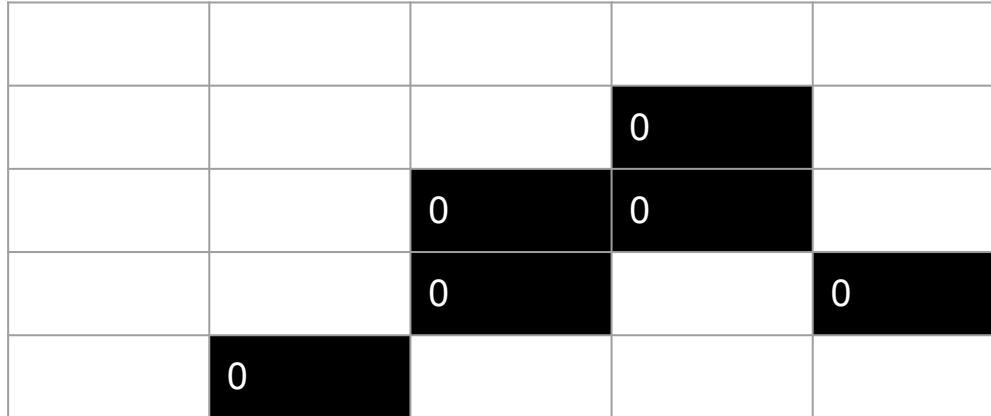
4-neighbors

0	255	0
255	255	255
0	255	0

8-neighbors

255	255	255
255	255	255
255	255	255

# Pixel Neighborhoods and Connectivity

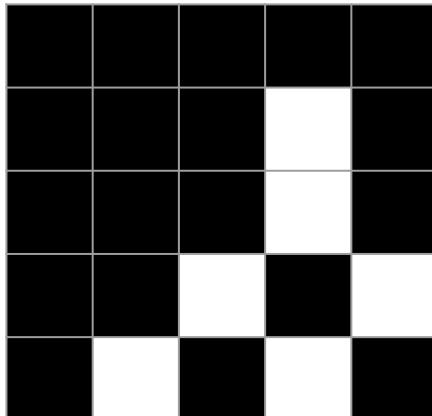


How many connected components?

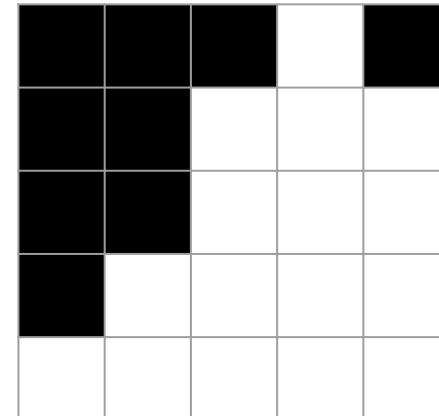
# Morphological Operations

- **Dilation** - expansion of the white connected-component by a filter.

$$A \oplus B = \{ x \mid (B)_x \cap A \neq \emptyset \}$$



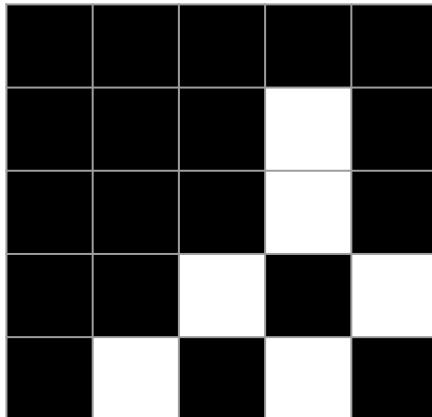
0	255	0
255	255	255
0	255	0



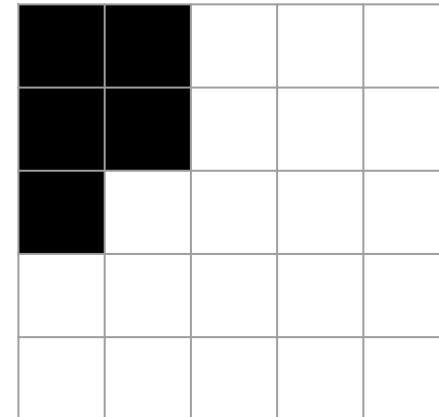
# Morphological Operations

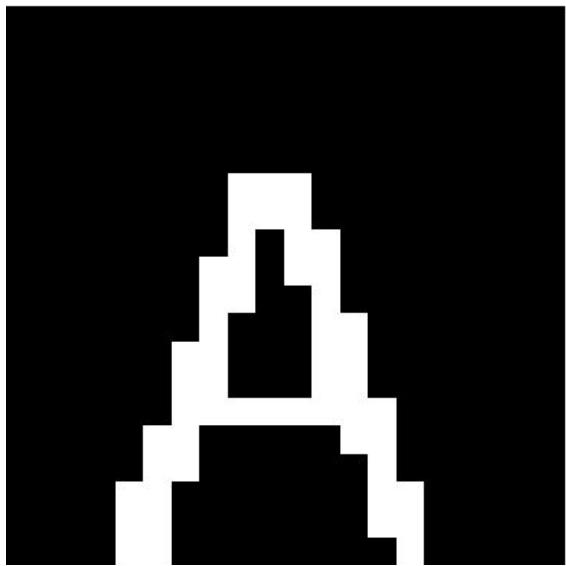
- **Dilation** - expansion of the white connected-component by a filter.

$$A \oplus B = \{ x \mid (B)_x \cap A \neq \emptyset \}$$

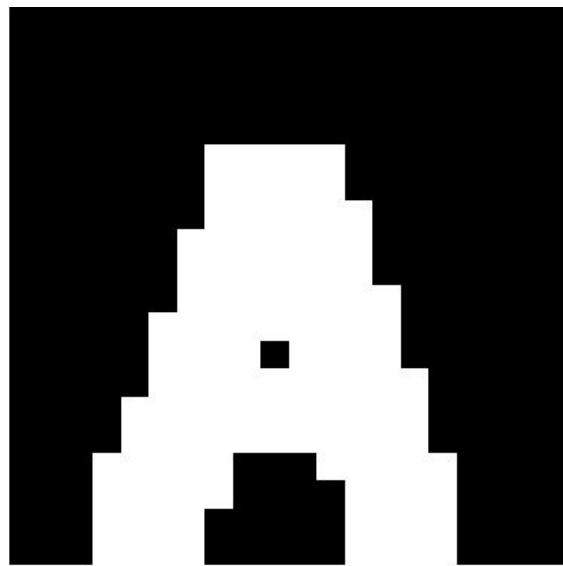


255	255	255
255	255	255
255	255	255





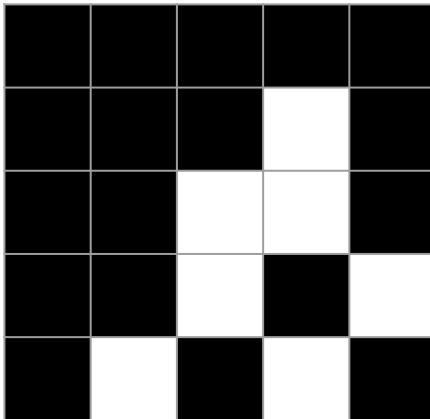
255	255	255
255	255	255
255	255	255



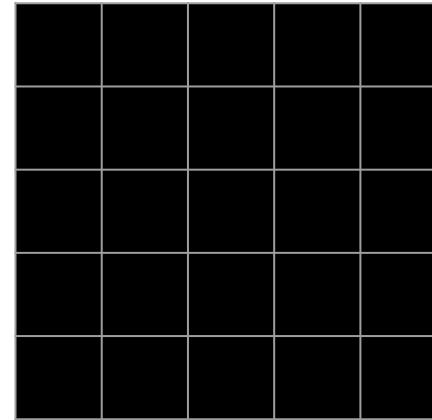
# Morphological Operations

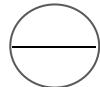
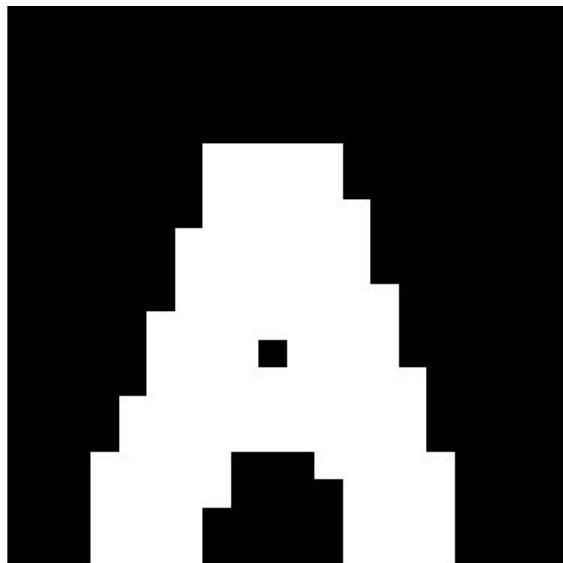
- **Erosion** - shrinking of the white connected-component by a filter.

$$A \ominus B = \{ x \mid (B)_x \subseteq A \}$$

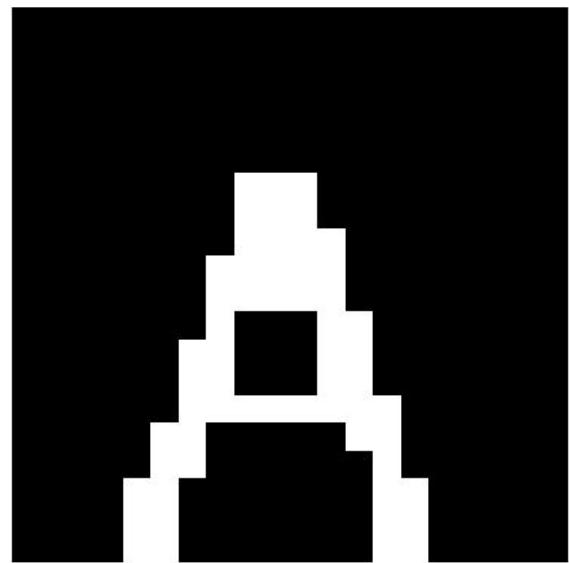


255	255	255
255	255	255
255	255	255

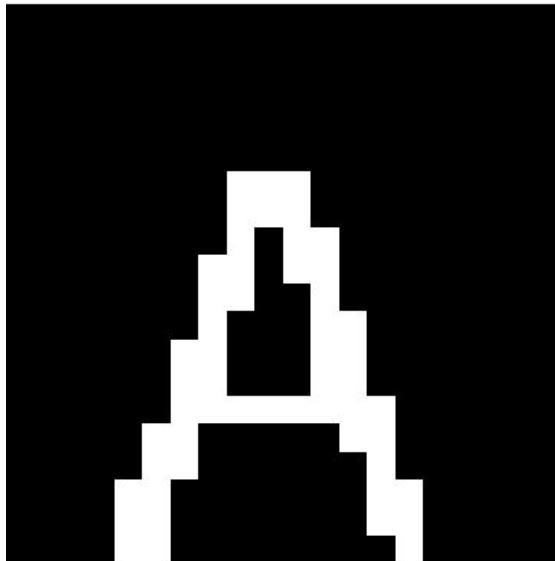




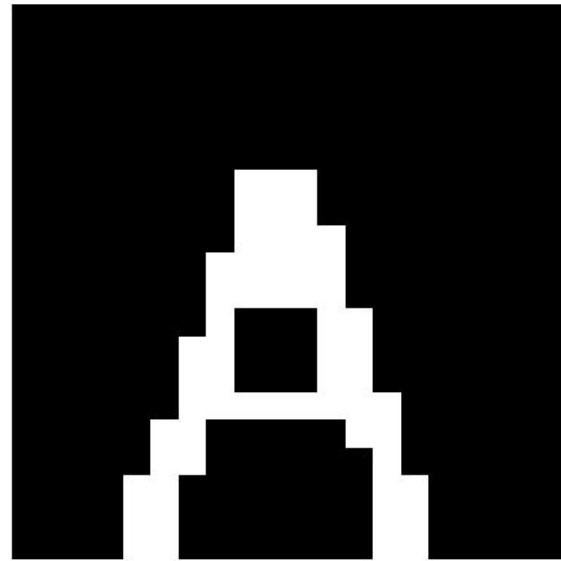
255	255	255
255	255	255
255	255	255



Before dilation



After erosion(dilation(im))

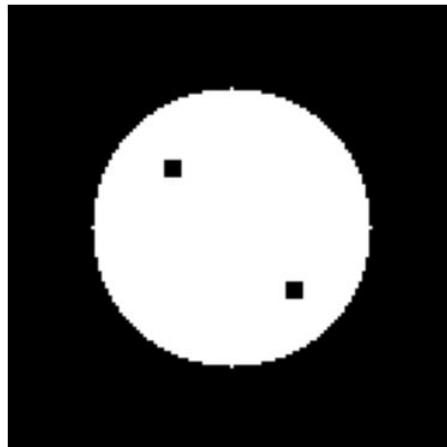


**Note – these are not opposites!  $\text{erosion}(\text{dilation}(\text{im})) \neq \text{im}$**

# Morphological Operations

- **Closing** - Dilation → Erosion  
Good for filling holes

Dilated

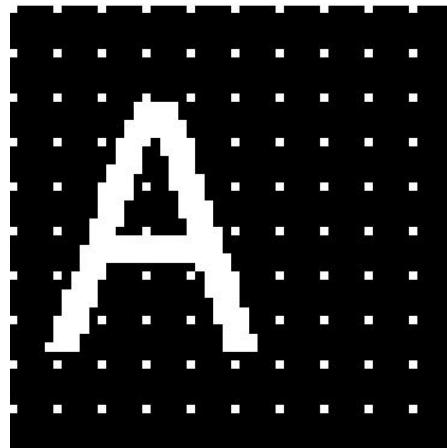


Eroded

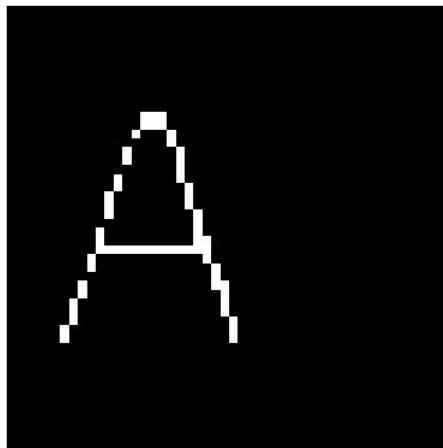


# Morphological Operations

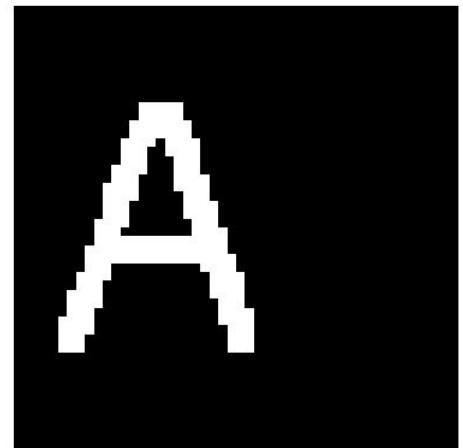
- **Opening** - Erosion → Dilation  
Good for removing noise



Eroded

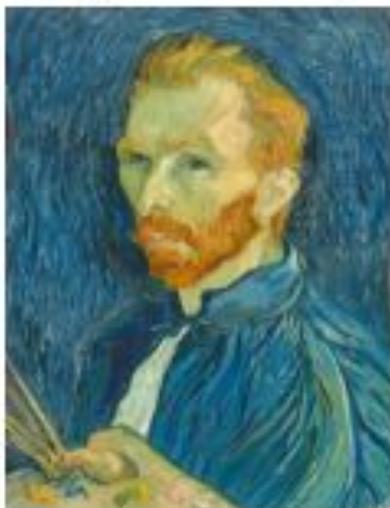


Dilated



# Image Resizing

Original image



im/4



im/8



Why does this happen?

# Image Resizing

Solution – blur before sub-sampling!

Original image



$G(im)/4$



$G(im)/8$



Note – Filter size should double for each  $\frac{1}{2}$  size reduction. Why?

# Image Resizing

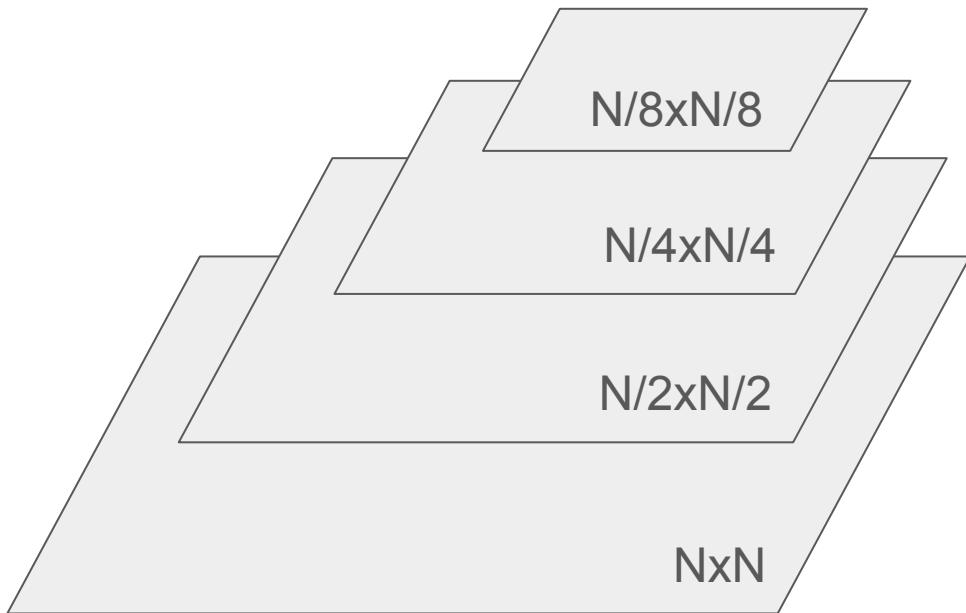
## Aliasing

- **Reduce:**
  - Blur
  - Sub-sample: select every n-th pixel in every n-th row.
- **Expand:**
  - Zero padding of every other pixel and row (e.g. for n=2: 1, 0, 2, 0, 3, 0,...)
  - Blur



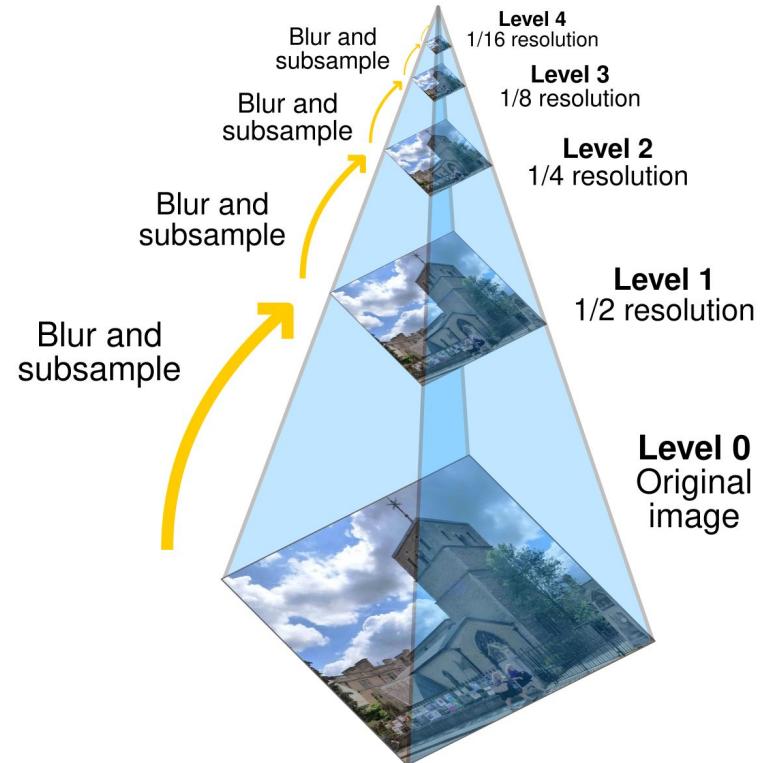
# Image Pyramids

- **Multi-scale template matching / object / feature detection:** Searching for patterns at different scales efficiently.
- **Image blending / seamless stitching:** combine images smoothly with Laplacian pyramids.
- **Efficient optical flow / motion estimation:** Coarse-to-fine pyramids let you estimate large motions at low resolution, then refine at higher resolution.
- **Compression:** Gaussian pyramids are used in wavelet-based or pyramid-based image compression.



# Gaussian Pyramid

Each level is **reduced** from the previous – gaussian blur → subsample.

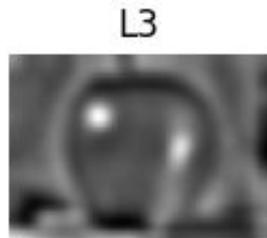
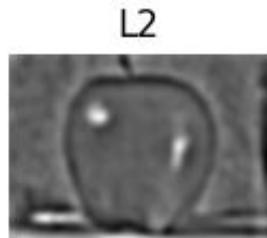
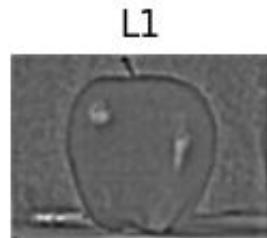
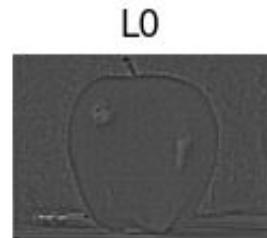
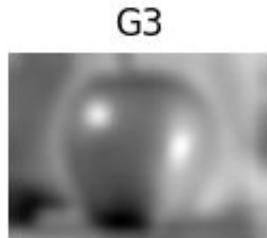
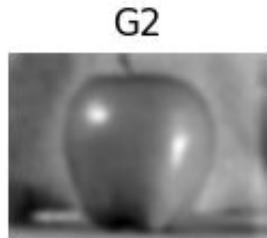
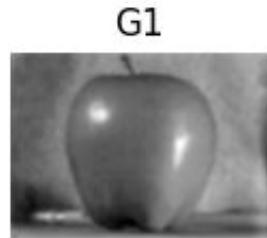
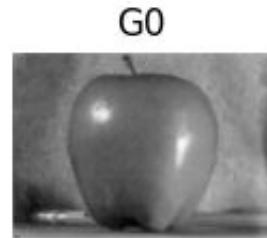


- Pyramid image from wikipedia

# Laplacian Pyramid

$$L_n = G_n$$

$$L_{n-1} = G_{n-1} - \text{Expand}(G_n)$$



# Laplacian Pyramid

$$L_n = G_n$$

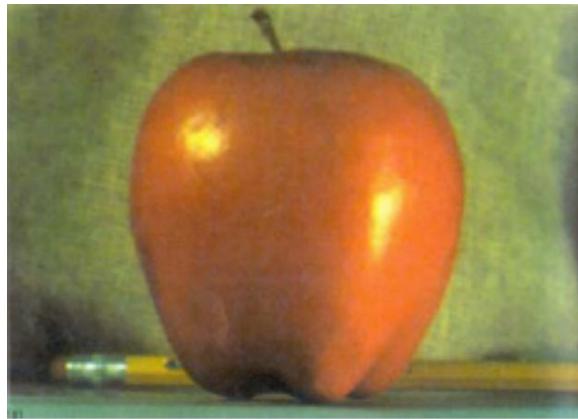
$$L_{n-1} = G_{n-1} - \text{Expand}(G_n)$$

$$\text{Expand}(L_n) + L_{n-1} = \text{Expand}(G_n) + G_{n-1} - \text{Expand}(G_n) = G_{n-1}$$

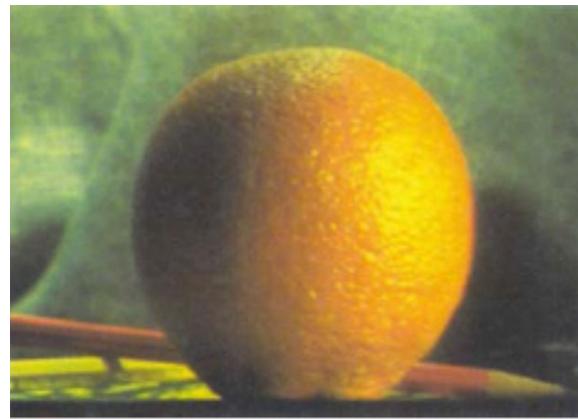
$$\sum_{i=0}^n L_i = G_0$$

- **Gaussian** Pyramid in the Fourier domain – Multiplication with a Gaussian kernel.
- **Laplacian** Pyramid in the Fourier domain – difference between two powers of Gaussian kernels.
- Sum of all laplacian layers = the original image.

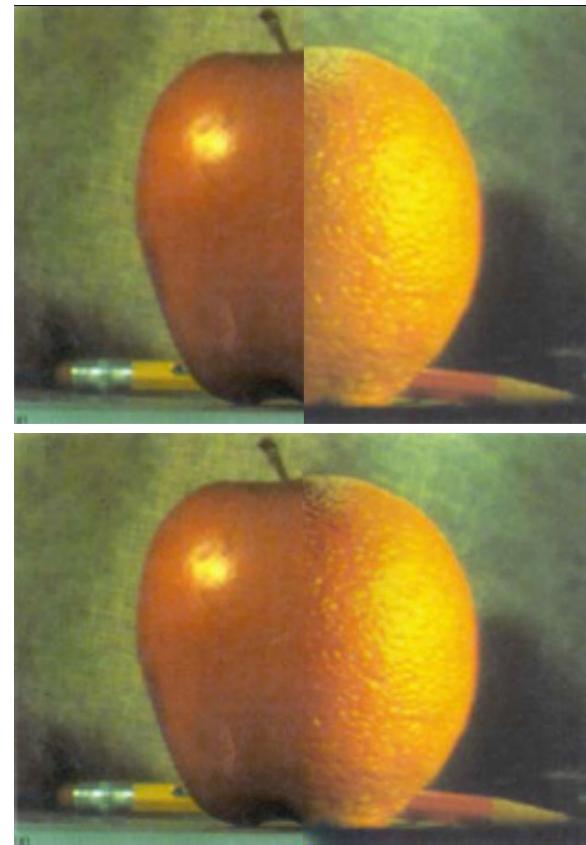
# Image Blending with Laplacian Pyramids



+

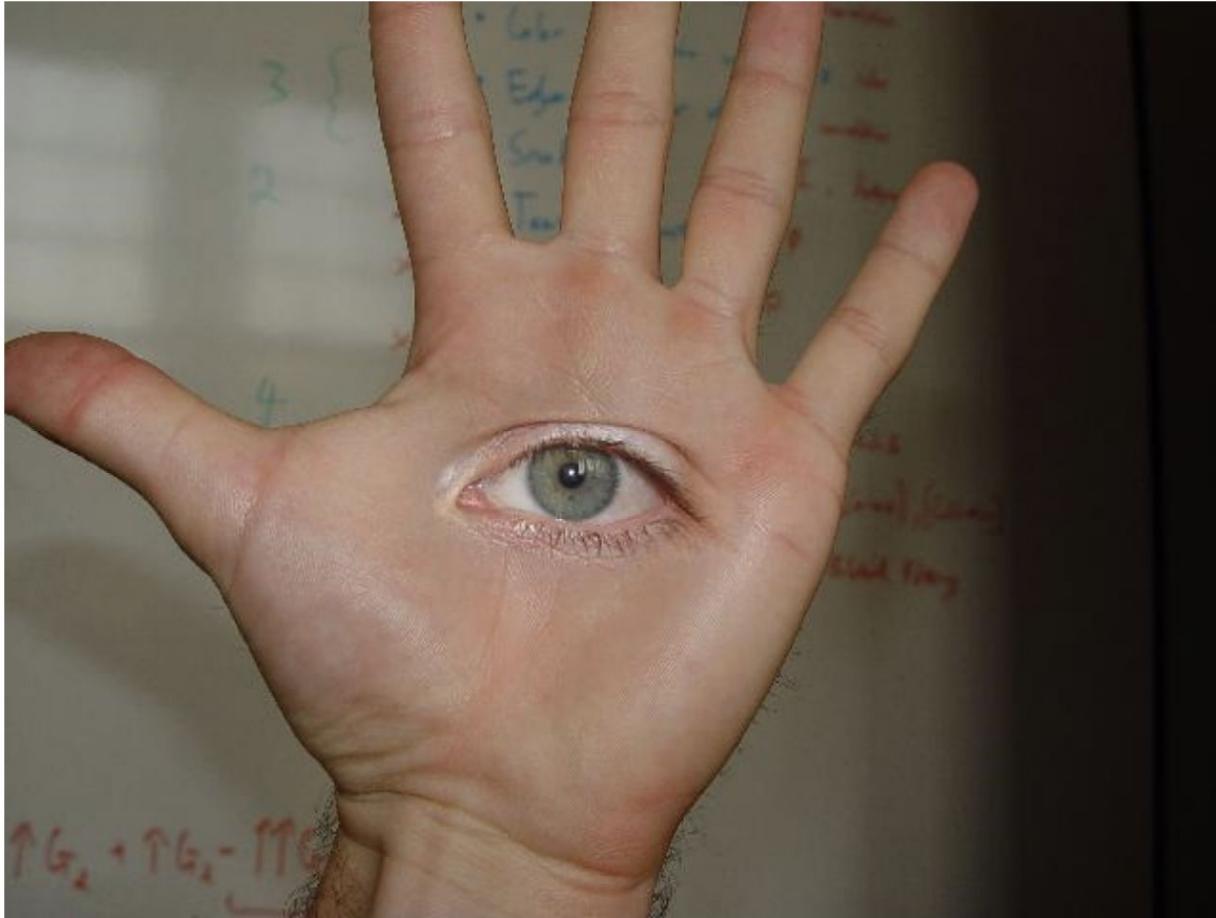


=



# Image Blending with Laplacian Pyramids

- **Input:** 2 images ( $im1$ ,  $im2$ ) and a binary mask ( $M$ )
- **Output:** a blended image based on the mask
- **Algorithm:**
  - Construct laplacian pyramids for each image:  $L1$ ,  $L2$ .
  - Construct a gaussian pyramid for the mask:  $Gm$
  - Construct another laplacian pyramid where for each level  $k$ :  
$$L3[k] = Gm[k]*L1[k] + (1-Gm[k])*L2[k]$$
  - Sum all levels in  $L3$  to get the blended image



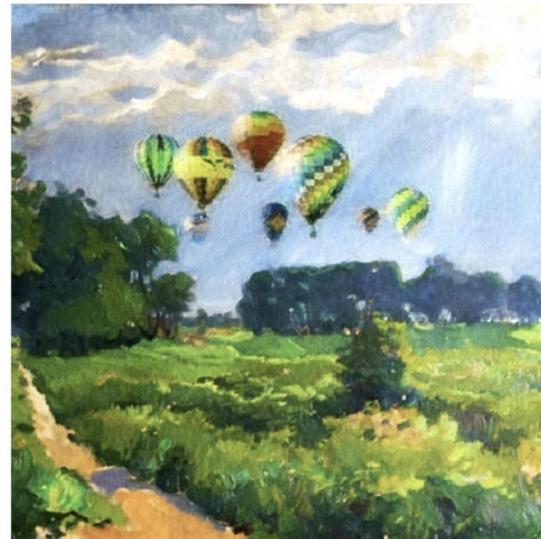
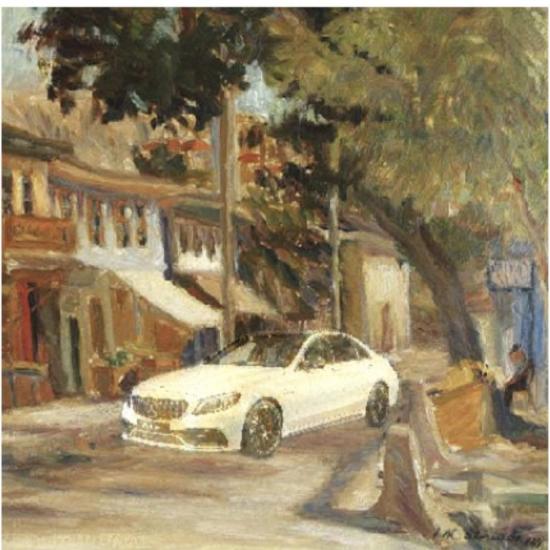
© prof. d'martin

# Newer blending methods

GP-GAN (Wu et al., ACMMM 2019)



Deep Image Blending (Zhang et al., WACV 2020)



# Blended Latent Diffusion (Avrahami et al., SIGGRAPH 2023)



Input Image



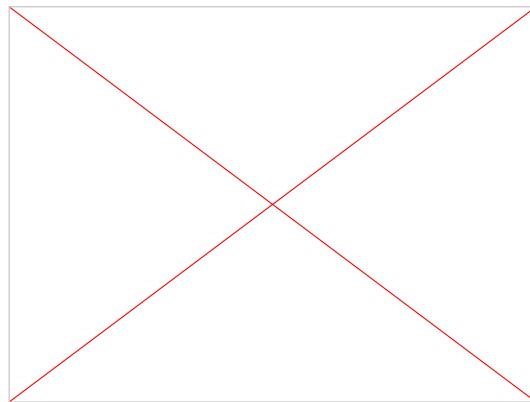
+

"A man with a  
+ yellow  
sweater"



Result

Split-then-Merge (Kara et al., Nov' 2025)



Input Prompt

GenCompositor (Yang et al., Sept' 2025)

